

Predictive_Modelling (1)

March 25, 2025

1 Predictive Modeling for Bengaluru Ride-Sharing Data

```
[4]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.metrics import classification_report, mean_squared_error

[5]: # Load and prepare data
df = pd.read_csv('bengaluru_ride_data.csv')

[6]: # Feature engineering
#df['DateTime'] = pd.to_datetime(df['Date'] + ' ' + df['Time'])
#df['Hour'] = df['DateTime'].dt.hour
#df['DayOfWeek'] = df['DateTime'].dt.day_name()
# Feature engineering
df['DateTime'] = pd.to_datetime(df['Date'] + ' ' + df['Time'])
df['Hour'] = df['DateTime'].dt.hour
df['DayOfWeek'] = df['DateTime'].dt.day_name()
df['Month'] = df['DateTime'].dt.month_name()
df['IsWeekend'] = df['DayOfWeek'].isin(['Saturday', 'Sunday']).astype(int)
df['IsPeakHour'] = (((df['Hour'] >= 7) & (df['Hour'] <= 10)) | ((df['Hour'] >= 17) & (df['Hour'] <= 20))).astype(int)

[7]: # Create target variables
df['IsCancelled'] = df['Booking Status'].isin(['Cancelled by Driver', 'Cancelled by Customer']).astype(int)
df['IsSuccessful'] = (df['Booking Status'] == 'Successful').astype(int)
```

2 Cancellation Prediction Model

```
[8]: # Prepare data for cancellation prediction
cancellation_features = ['Vehicle Type', 'Hour', 'DayOfWeek', 'Month', 'IsWeekend',
                        'IsPeakHour', 'Pickup Location', 'Drop Location']
X_cancel = df[cancellation_features]
y_cancel = df['IsCancelled']
```

```
[9]: # Preprocessing pipeline
categorical_features = ['Vehicle Type', 'DayOfWeek', 'Month', 'Pickup Location', 'Drop Location']
numeric_features = ['Hour', 'IsWeekend', 'IsPeakHour']

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ])

```

```
[9]:
```

```
[10]: # Preprocessing pipeline
categorical_features = ['Vehicle Type', 'DayOfWeek', 'Month', 'Pickup Location', 'Drop Location']
numeric_features = ['Hour', 'IsWeekend', 'IsPeakHour']

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ])

```

```
[11]: # Split data
X_train, X_test, y_train, y_test = train_test_split(
    X_cancel, y_cancel, test_size=0.2, random_state=42)
```

```
[12]: # Model training (using XGBoost for better performance)
from xgboost import XGBClassifier
```

```
[13]: cancel_model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('classifier', XGBClassifier(random_state=42))
])

cancel_model.fit(X_train, y_train)
```

```
[13]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num', StandardScaler(),
                                                         ['Hour', 'IsWeekend',
                                                          'IsPeakHour']),
                                                         ('cat',
                                                          OneHotEncoder(handle_unknown='ignore'),
                                                         ['Vehicle Type', 'DayOfWeek',
                                                          'Month', 'Pickup Location',
                                                          'Drop Location'])])),
                    ('classifier',
                     XGBClassifier(base_score=None, booster=None, callbacks=None,
                                   colsample_bylevel=None, colsample_bynode=None,
                                   feature_types=None, gamma=None, grow_policy=None,
                                   importance_type=None,
                                   interaction_constraints=None, learning_rate=None,
                                   max_bin=None, max_cat_threshold=None,
                                   max_cat_to_onehot=None, max_delta_step=None,
                                   max_depth=None, max_leaves=None,
                                   min_child_weight=None, missing=None,
                                   monotone_constraints=None, multi_strategy=None,
                                   n_estimators=None, n_jobs=None,
                                   num_parallel_tree=None, random_state=42, ...))])
```

```
[14]: # Evaluate
y_pred = cancel_model.predict(X_test)
print("Cancellation Prediction Performance:")
print(classification_report(y_test, y_pred))
```

```
Cancellation Prediction Performance:
              precision    recall  f1-score   support

    0               0.70         0.96         0.81         2459
    1               0.28         0.04         0.07         1043

 accuracy               0.68         0.68         0.68         3502
 macro avg              0.49         0.50         0.44         3502
 weighted avg           0.58         0.68         0.59         3502
```

3 Demand Forecasting Model

```
[25]: # Prepare time series data for demand forecasting
demand_df = df.groupby(['Date', 'Hour', 'Vehicle Type']).size().
    .reset_index(name='Demand')
demand_df['Date'] = pd.to_datetime(demand_df['Date'])
demand_df['DayOfWeek'] = demand_df['Date'].dt.day_name()
```

```
demand_df['IsWeekend'] = demand_df['DayOfWeek'].isin(['Saturday', 'Sunday']).
↳ astype(int)
```

```
[26]: # Feature engineering for demand
demand_df = demand_df.sort_values(['Vehicle Type', 'Date', 'Hour'])
```

```
[27]: # Calculate lag features by group
def calculate_lags(group):
    group = group.sort_values('Date')
    group['Demand_Lag_1'] = group['Demand'].shift(1)
    group['Demand_Lag_24'] = group['Demand'].shift(24) # Same hour previous day
    group['Demand_Lag_168'] = group['Demand'].shift(168) # Same hour same day,
↳ previous week
    group['Rolling_Avg_24'] = group['Demand'].rolling(24, min_periods=1).mean()
    return group

demand_df = demand_df.groupby('Vehicle Type').apply(calculate_lags).
↳ reset_index(drop=True)
```

<ipython-input-27-44772df39df5>:10: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns. This behavior is deprecated, and in a future version of pandas the grouping columns will be excluded from the operation. Either pass `include_groups=False` to exclude the groupings or explicitly select the grouping columns after groupby to silence this warning.

```
demand_df = demand_df.groupby('Vehicle
Type').apply(calculate_lags).reset_index(drop=True)
```

```
[28]: # Drop rows where all lag features are NA (first week of data)
demand_df = demand_df.dropna(subset=['Demand_Lag_1', 'Demand_Lag_24',
↳ 'Demand_Lag_168'], how='all')
```

```
[29]: # Check if we have data remaining
if len(demand_df) == 0:
    raise ValueError("No data remaining after processing lag features. Check,
↳ your input data.")
```

```
[30]: # Prepare features and target
X_demand = demand_df[['Hour', 'Vehicle Type', 'DayOfWeek', 'IsWeekend',
    'Demand_Lag_1', 'Demand_Lag_24', 'Demand_Lag_168',
↳ 'Rolling_Avg_24']]
y_demand = demand_df['Demand']
```

```
[31]: # Preprocessing
demand_preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Hour', 'IsWeekend', 'Demand_Lag_1',
↳ 'Demand_Lag_24', 'Demand_Lag_168', 'Rolling_Avg_24']),
```

```

        ('cat', OneHotEncoder(handle_unknown='ignore'), ['Vehicle Type',
↪ 'DayOfWeek'])
    ])

```

```

[32]: # For time series data, we should do a temporal split rather than random split
split_point = int(0.7 * len(demand_df))
X_train_d = X_demand.iloc[:split_point]
X_test_d = X_demand.iloc[split_point:]
y_train_d = y_demand.iloc[:split_point]
y_test_d = y_demand.iloc[split_point:]

```

```

[33]: # Verify we have data in both sets
if len(X_train_d) == 0 or len(X_test_d) == 0:
    raise ValueError("Train or test set is empty after splitting. Adjust split_
↪ ratio or check data.")

```

```

[34]: # Model training (using Random Forest for time series)
from sklearn.ensemble import RandomForestRegressor

demand_model = Pipeline(steps=[
    ('preprocessor', demand_preprocessor),
    ('regressor', RandomForestRegressor(n_estimators=100, random_state=42))
])

demand_model.fit(X_train_d, y_train_d)

```

```

[34]: Pipeline(steps=[('preprocessor',
                        ColumnTransformer(transformers=[('num', StandardScaler(),
                                                         ['Hour', 'IsWeekend',
                                                         'Demand_Lag_1',
                                                         'Demand_Lag_24',
                                                         'Demand_Lag_168',
                                                         'Rolling_Avg_24']),
                                                         ('cat',
                                                         OneHotEncoder(handle_unknown='ignore'),
                                                         ['Vehicle Type',
                                                         'DayOfWeek'])])),
                        ('regressor', RandomForestRegressor(random_state=42))])

```

```

[36]: from sklearn.metrics import mean_absolute_error # Import the missing function

```

```

[37]: # Evaluate
y_pred_d = demand_model.predict(X_test_d)
mse = mean_squared_error(y_test_d, y_pred_d)
print(f"\nDemand Forecasting Performance:")
print(f"MSE: {mse:.2f}")
print(f"RMSE: {np.sqrt(mse):.2f}")

```

```
print(f"Mean Absolute Error: {mean_absolute_error(y_test_d, y_pred_d):.2f}")
```

Demand Forecasting Performance:

MSE: 2.69

RMSE: 1.64

Mean Absolute Error: 1.23

4 Fare Prediction Model

```
[38]: # Prepare data for fare prediction
fare_df = df[df['Booking Status'] == 'Successful'].copy()
fare_features = ['Vehicle Type', 'Ride Distance (km)', 'Hour', 'DayOfWeek',
                 'Month', 'IsWeekend', 'IsPeakHour', 'Pickup Location', 'Drop_
                 Location']
X_fare = fare_df[fare_features]
y_fare = fare_df['Booking Value ( )']
```

```
[39]: # Preprocessing
fare_preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), ['Ride Distance (km)', 'Hour', 'IsWeekend',
        IsPeakHour']),
        ('cat', OneHotEncoder(handle_unknown='ignore'), ['Vehicle Type',
        DayOfWeek', 'Month', 'Pickup Location', 'Drop Location'])
    ])

```

```
[40]: # Split data
X_train_f, X_test_f, y_train_f, y_test_f = train_test_split(
    X_fare, y_fare, test_size=0.2, random_state=42)
```

```
[41]: # Model training (using Gradient Boosting for regression)
from sklearn.ensemble import GradientBoostingRegressor

fare_model = Pipeline(steps=[
    ('preprocessor', fare_preprocessor),
    ('regressor', GradientBoostingRegressor(random_state=42))
])

fare_model.fit(X_train_f, y_train_f)
```

```
[41]: Pipeline(steps=[('preprocessor',
    ColumnTransformer(transformers=[('num', StandardScaler(),
    ['Ride Distance (km)', 'Hour',
    'IsWeekend', 'IsPeakHour']),
    ('cat',
```

```
OneHotEncoder(handle_unknown='ignore'),
                ['Vehicle Type', 'DayOfWeek',
                 'Month', 'Pickup Location',
                 'Drop Location']]))),
                ('regressor', GradientBoostingRegressor(random_state=42)))]
```

```
[42]: # Evaluate
y_pred_f = fare_model.predict(X_test_f)
mse = mean_squared_error(y_test_f, y_pred_f)
print(f"\nFare Prediction MSE: {mse:.2f}")
print(f"RMSE: {np.sqrt(mse):.2f}")
```

Fare Prediction MSE: 106990.86
RMSE: 327.09


5 Model Deployment and Feature Importance

```
[43]: # Feature importance for cancellation model
import matplotlib.pyplot as plt

# Get feature names after one-hot encoding
cancel_features = (numeric_features +
                   list(cancel_model.named_steps['preprocessor']
                       .named_transformers_['cat']
                       .get_feature_names_out(categorical_features)))
```

```
[44]: # Get feature importances
importances = cancel_model.named_steps['classifier'].feature_importances_
```

```
[45]: # Plot
plt.figure(figsize=(12, 8))
pd.Series(importances, index=cancel_features).sort_values(ascending=True).
    ↪tail(20).plot.barh()
plt.title('Top 20 Features for Cancellation Prediction')
plt.show()
```



```
[46]: # Save models for deployment
import joblib
joblib.dump(cancel_model, 'cancellation_model.pkl')
joblib.dump(demand_model, 'demand_model.pkl')
joblib.dump(fare_model, 'fare_model.pkl')
```

```
[46]: ['fare_model.pkl']
```

The predictive modeling project focused on three key areas for Bengaluru ride-sharing data:

Cancellation Prediction:

Goal: Predict whether a booking will be canceled (binary classification).

Results: Achieved 68% accuracy, with high recall (96%) for non-cancellations but lower performance for cancellations (precision: 28%).

Key Features: Hour, peak hours, and vehicle type were most influential.

Demand Forecasting:

Goal: Forecast ride demand by hour and vehicle type (time series regression).

Results: Low error rates (MAE: 1.23, RMSE: 1.64), indicating reliable predictions.

Approach: Used lag features (e.g., previous day/hour demand) and rolling averages.

Fare Prediction:

Goal: Estimate ride fares for successful bookings (regression).

Results: Higher error (RMSE: 327.09), likely due to fare variability. Suggests room for improvement (e.g., more features or model tuning).

Next Steps:

Address class imbalance in cancellation predictions (e.g., resampling or cost-sensitive learning).

Refine fare model with additional features (e.g., traffic, weather) or alternative algorithms.

Deploy models (`cancellation_model.pkl`, `demand_model.pkl`, `fare_model.pkl`) for real-time insight

[]: